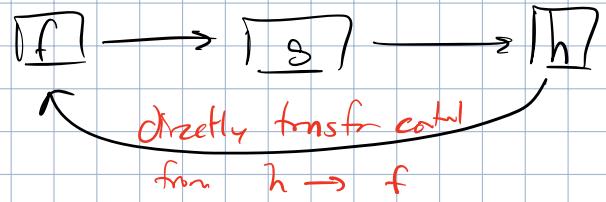


NON-LOCAL TRANSFER



In single func., multi-level exits are easy \rightarrow goto label

\hookrightarrow Problem: labels are scoped to function

Another problem to solve: multiple possible return results!

Traditional Approaches to Multiple Returns

① Return code: indicates normal/exceptional return

$\hookrightarrow \text{printf}() \rightarrow 0 = \text{good}, -1 = \text{exception}$

② Status flags: global variables that indicate exception/normal

③ Fixup: pass fixup func. to call when exception occurs

④ Return union: combine code w/ object

Go { $\begin{matrix} \downarrow \text{stat} \\ \text{ok}, \end{matrix}$, $\begin{matrix} \uparrow \text{val} \\ \text{val} := \text{foo}() \end{matrix}$ }
 \hookrightarrow requires checking 'ok' before 'val'

Drawbacks:

1. Checks can be optional
2. Mixing return w/ exceptions
3. Multiple exception types \Rightarrow bloat in return types/codes
4. Opaque code/statistics
5. Increasing param # w/ fixup

Dynamic Multi-level Exit

① Set a global label var

② B4 func call, set label var

③ Throw exception \rightarrow jump to global label

④ On any exit/return/exception handling, reset to prev label

Skips stack scopes! Will clean up stacks implicitly \rightarrow pseudo multi-level exit

- \rightarrow 1. Direct control flow to label (updating PC)
 2. Go to transfer point

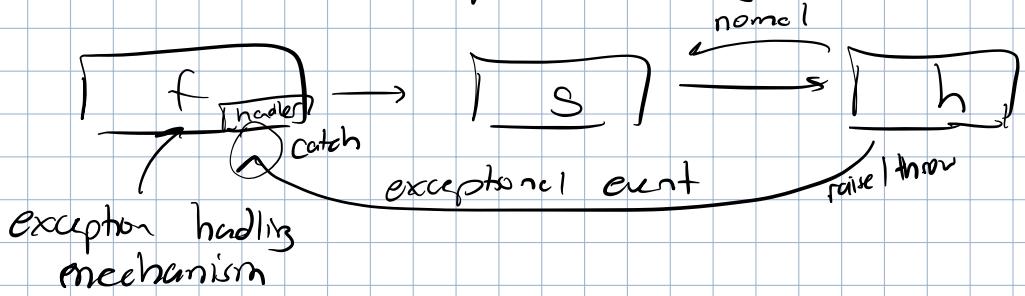
Dynamic: you can never know what global label is set to from reading code!

In C/C++: jmp-but → label
setjmp → sets global labl
longjmp → go to label

Problem: too general

Exception Handling

Non-local transfer ≡ exception handling



Types of execution:

- ① Source: execution (program w/ own stack) that raises an exception
 - ② Faults: n changes control flow (e.g. whatever catches)
- If ① = ② ⇒ local exception. O.W.: non-local exception.

Non-local exceptions can be concurrent, or not!

Propagation: direct exception from source → faults

↳ Order of catchers is most recent func. w/ com catch.

Once exception reaches handler:

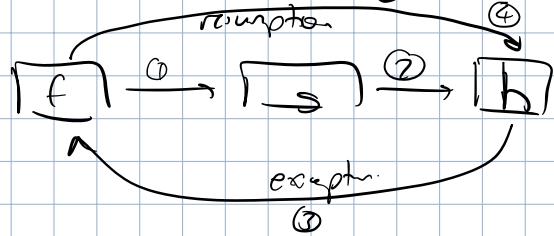
- ① Return → handler has handled exception
- ② Reraise (often to add extra info)
- ③ Raise a new (private exception)

In C++:

```
Success block { try {  
    ... } catch (exceptionType) {  
    ... }
```

Code here can throw an exception.

Resumption: Control flow goes back to raise



Q: What happens to stack object if exce. thrown?

↳ Terminate block: call others & code in 'finally'

try {

- -

{ catch () { - - } will always be exec.
{ finally { - - }}

Note: we don't directly jump to routine that handles exception → stack unwinding must happen.

Implementation

Approach ①: Use labels

1. Assoc. global label / exception type
2. Set label on guarded block entry
3. Reset global label of catch / exit of block

Con: lot of ops. if lots of try & catches. Exceptions not raised too often

Approach ②: Lookup table

- ① Store catch & other info per block.
 - ② Exc. raised → do linear walk in LUT
- } amortizes cost on exception raising?

Static / Dynamic Call / Return

- Static / dynamic call: routine name is known @ compilation (static) or not (dynamic)

- Static / dynamic return: routine ret. location known

call

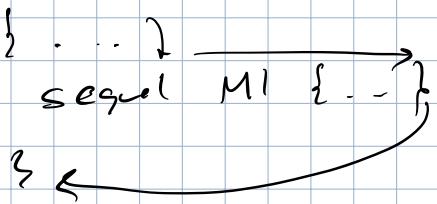
return static
static sequel

dynamic reg routine

dynamic
terminator exception
passing routine into func.
/ resumption

Sequel: routine w/ no return value.

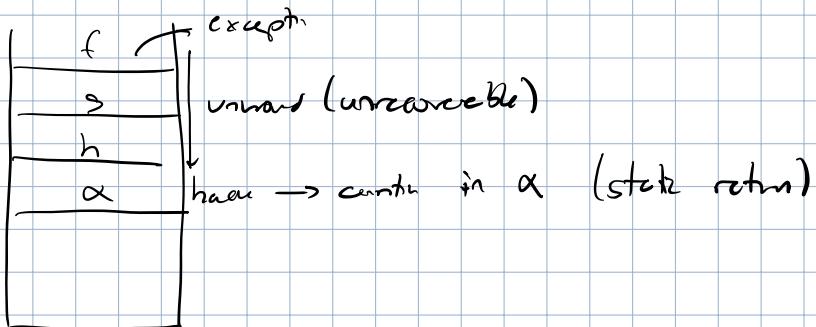
↳ Control returns to end of block where sequel is located



* Often used for non-recoverable event (e.g. FileException, StackOverflow) *

Not good for modularize

Termination: dynamic raise (handler is unknown), static return b/c we end handle \leftarrow stack is unknown.



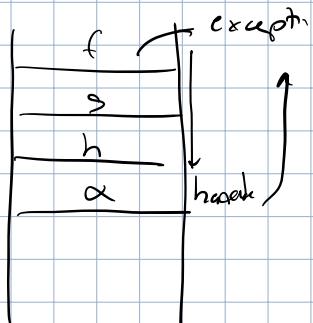
In C++:

```
foo () {  
    ...  
    throw E();  
}  
bar () {  
    ...  
    try {  
        ...  
        foo();  
    } catch (E) {  
        // specify type  
        // if type is  
        // catch error  
    }  
}
```

Key note: if you don't want to handle throws \rightarrow no need to make catch \leftarrow goes down stack

try...retry: restart try block after exception handled

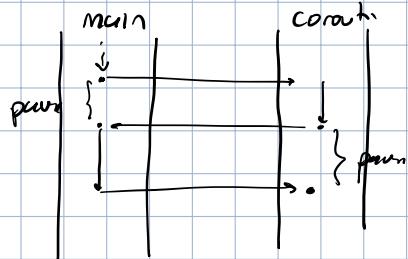
Resumption: dynamic raise b/c handler is unknown but stack is not unknown. At end of handler, don't know where you will resume \rightarrow dynamic return.



COROUTINES

Defn: routine that can be passed & reused.
↳ Why? Has its own stack! Same stack.

Not concurrency \leftarrow not doing simultaneous comp.



Semi-conductors

Struktur:

```

- Coroutine <none> {
    private members: ← comm.
    void main() {
        --- ⇒ suspend in here
    }
}

publiz:
    von foo() { - resume() -- }

```

1st closure: starts new stack for coroutine.

Suspend: reactivates last resume () call

Coroutines can be deleted b/f done

Since `Suspend()` will save states on stack \rightarrow no need for flag variable!

Error handling

① Defining exceptions:

```
- Coroutine <--> {  
    private:  
        - - -  
    public:  
        ✓ class exception  
        - Event <ExceptionName> { .. }  
        :  
        :  
    }{
```

⑦ Raising exceptions:

- Throw ExceptionName(); or - Pressure ExceptionName() [- At construction] Ls Go with this

-Pressure Exemption (L) [- At certain] Ls Go with this

In µ++: exception caught by lowest level class, not parent class

③ Handle exception:

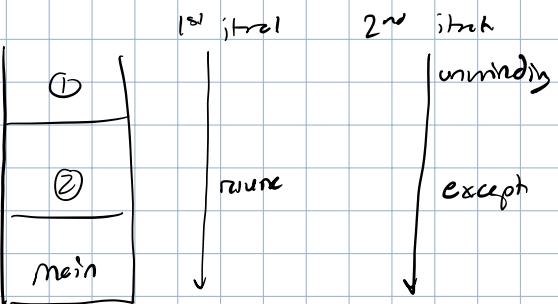
```
void main () {  
    try {  
        - Enable {  
            . . .  
        } - Catch Resume (...) {  
            . . .  
        } catch (...) {  
            . . .  
    }  
}
```

Allows exceptions returned / thrown at routine
to be handled

Resume: no stack unwinding

Exception: stack unwinding

Q: What happens if -Resume but no -CatchResume?



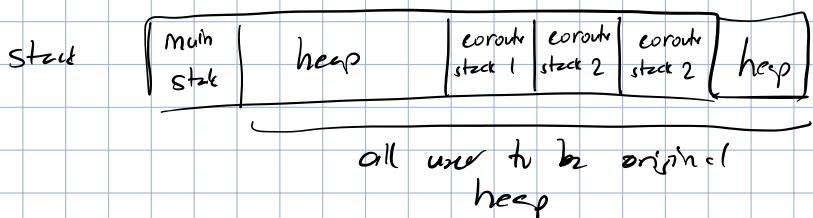
Multiple exceptions at coroutine → queued & in FIFO

Can always disable exception handling via -Disable {} - ?

Starter: program that does first resume

↳ Semicoroutine: only starter calls resume (not multiple coroutines calling resume())

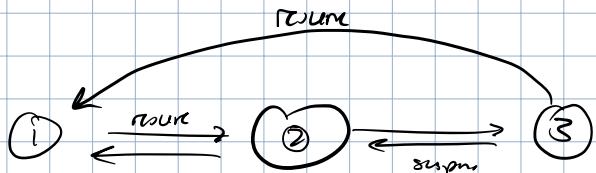
Memory management



Stacks small in coroutines → Don't allow big objects

Full coroutines

Resume cycle:



`resume ()`:
 ① Pause curr. routine
 ② Activate curr. object

`suspend ()`:
 ① Pause curr corout.
 ② Activate last resurer

Constructing full coroutine:

A: Coroutines need references to each other:

- Coroutine Example {

Example next; ← returns set refence
 private:

```
void main {
    cout << "ID: " << id << endl;
    next. call();
}
```

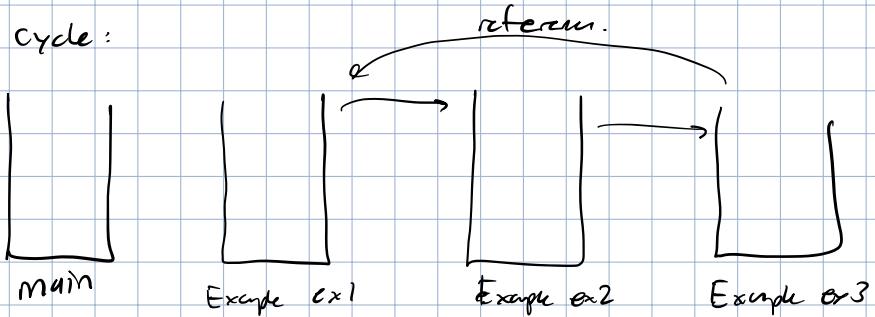
public:

```
Example (Example next) : next(next) {};
```

```
void call {
    resume();
}
```

If cycle of references, might need special function to make references

B: Start cycle:



1. main. call() ex1. call():

resume() → inactive main
 activate "this" → ex1

2. "ID: 1"

3. call ex2. call()

resume() → inactive ex1
 activate "this" → ex2

4. "ID: 2"

5. call ex3. call()

6. "ID: 3"

:

C: End cycle

Get back to start, but not always the last runner

Ending coroutine → start

Coroutine in modern languages

Coroutine implementations

Stackless

- Use caller stack
- Fixed size state

Stackful

- Separate stack } MC++
- Fixed size stack } MC++

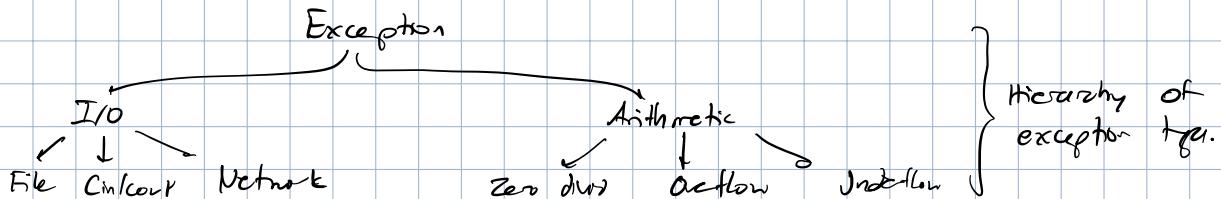
Python: yield & next
↓
suspense resource

JS: yield & ...next()

C++: has API but not good

More Exceptions

Derived Exception Type



Hierarchy enables programer to handle exceptions at diff. granularities

↳ Less coupling for specific exceptions

To handle specific & general case:

```
try {  
    ...  
} catch (SpecificExc & e) {  
    ...  
} catch (GeneralExc & e) {  
    ...  
}
```

Note: catch by reference ⇒ dynamic casting, re-resumes

Catch - Any

Purpose: catch any sort of exception

Soln:

① Catch base exception class

② Syntax:

C++: } ... catch (...) { .. } \Rightarrow Applies to -CatchResume

③ Finally (in Java)

try { .. } catch (...) { .. } finally { .. } \swarrow Always executed

Exception Parameters

Goal: pass info from raise point \rightarrow handle

Soln:

1. Define struct for exception

2. At raise: initialize params

3. Catch by reference \rightarrow use info

L For resumption = change info in exception \rightarrow -Resume

Exception List

Goal: specify which routines are possibly throwing an exception \Rightarrow static checking

L Ex:// int f() throw (E, B, ..) { .. }

Java uses a static check, C++ uses dynamic runtime checks

Issue w/ exception lists:

1. Cannot know all exceptions from routine

2. Fixed set of exception types

Destructors

Most times, others will not throw exception. Can force it to throw exceptions

L class A {

~A() noexcept(false) { .. }

Multiple Exceptions

Not usually an issue unless other raises exception \Rightarrow undefined behavior

Multiple exceptions can be thrown, but only 1 can propagate

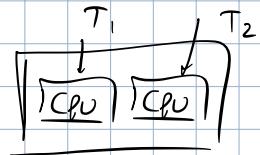
CONCURRENCY

Thread: indep. sequential exec. of prog w/ own state

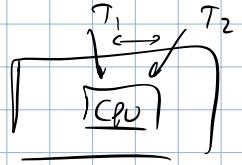
Process: has own threads & state into (like coroutines). Can have multiple threads.

Task: light-weight process, smaller & lighter

Parallel execution:



Concurrent execution:



switch so fast that it looks like parallel execution

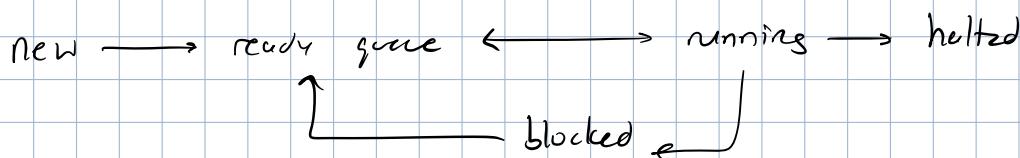
Why?

- ① Problem divided is easier to solve
- ② Speedups

Difficulty?

- ① Understanding coordination, simultaneous exec.
- ② Problem specification: how to break up problem, comm.
- ③ Debug: non-deterministic order

Execution States



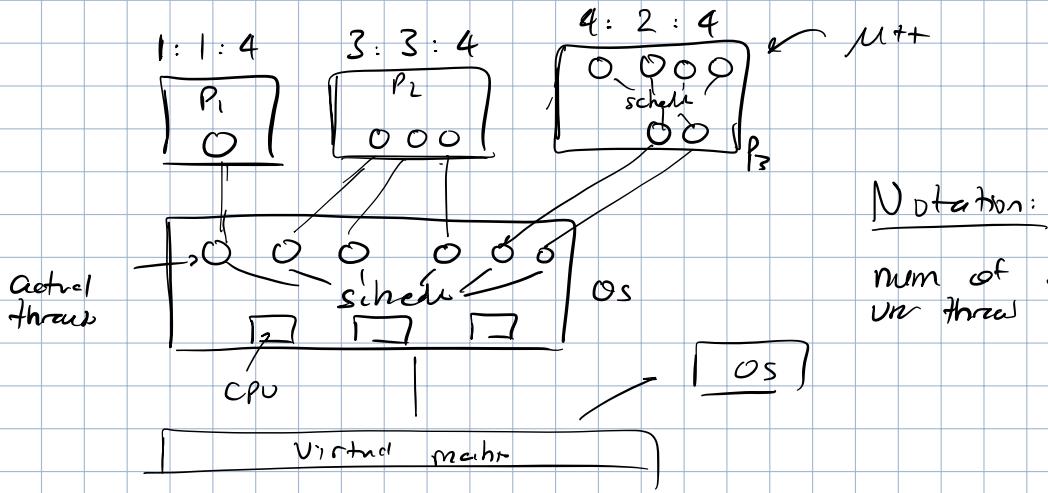
Scheduler determines who should run \rightarrow cannot skip ready queue
Operations that are non-atomic can switch from running state midway

↳ Can lead to data loss

Ex:// 3 threads running:

```
int global = 0  
for (int i = 0; i < N; i++) {  
    global += i  
}  
↳ Threads can switch halfway
```

Threading model



Notation:

num of user threads : num of actual threads : num of CPUs

Concurrent Systems

Types of concurrent systems:

① Automatically converts sequential \rightarrow concurrent prog.

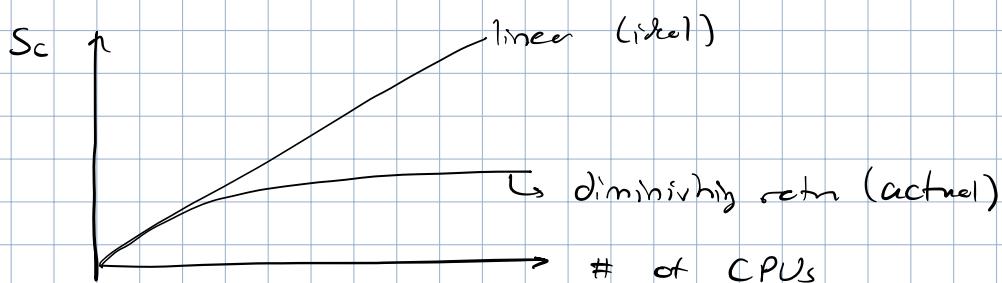
② Provide hints/annotation to sequential \rightarrow compiler converts to concurrent

③ Explicit concurrency constructs

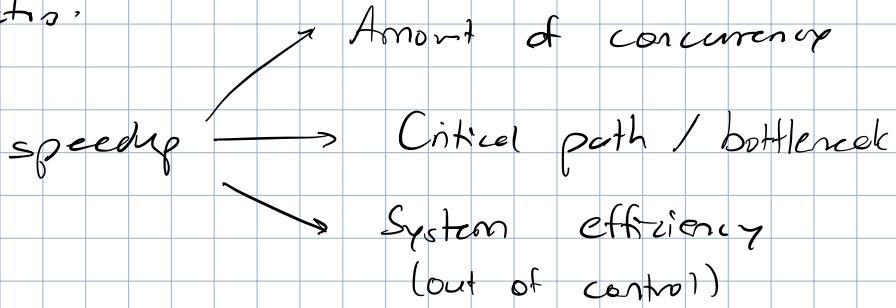
At some point, you need to use explicit concurrency

Speedup

$$S_c = \frac{T_1}{T_c} \quad \begin{matrix} \leftarrow \text{time in seq. pros} \\ \leftarrow \text{time in concn. pros.} \end{matrix}$$



Speedup factors:



Andahl's Law:

Speedup of C CPUs w/ prog taking P_i time in concurrency is:

$$S_C = \frac{1}{\underbrace{1-p}_{\text{seq.}} + p/c} \quad (\text{assuming } T_1 = 1, T_c = sq + \text{concn.})$$

Ex:// Seq. takes 10 seconds on 1 CPU. 2.5 sec for 4 CPUs w/ 100% concurrency

$$S_4 = \frac{1}{1 - \frac{1}{4} + \frac{1}{4}} = \frac{4}{4}$$

$p = 100\%$

Ex:// 80% concurrency:

$$\begin{aligned} S_4 &= \frac{1}{1 - 0.8 + \frac{0.8}{4}} \\ &= \frac{1}{0.2 + 0.2} \\ &= 2.5 \text{ times speedup} \end{aligned} \quad \begin{matrix} \text{Hyper decrease b/c} \\ \text{of sequential code} \end{matrix}$$

To get real speedup \rightarrow minimize time in sequential part.
& reduce bottlenecks

Thread Creation

Any concurrency construct requires:

1. Creation

2. Synchronization

3. Communication

① COBEGIN / COEND

#include <w Cobegin.h>

```
int main(--) {
```

```
; COBEGIN
```

```
BEGIN --- END → Each BEGIN in theory has own  
BEGIN --- END thread  
BEGIN --- END
```

```
; COEND
```

→ Wait for all threads to finish

```
}
```

Implicit concurrency

② START / WAIT

```
#include <vCobegin.h>
```

```
int main(--) {
```

```
auto ptr1 = START(func, params)
```

```
auto ptr2 = START(--)
```

```
;
```

```
;
```

i = WAIT(ptr1) → Unit until thread finished
& return the value of func

```
}
```

cofor (i, start, end, { - })
 Runs threads instead start
 → end - 1 that runs { - }
 can access thread i/x using i

Thread state
 running func w/
 params

Explicit concurrency

START / WAIT can simulate COBEGIN / COEND, not v.v.

③ Thread object

```
- Task <name> {
  void main() { ... } // Function runs if
}; // thread started
```

Initialize:

```
{ ← COBEGIN
```

```
;
```

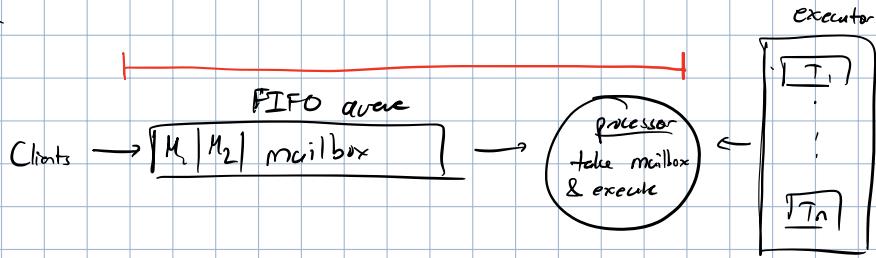
```
<name> obj ; OR
```

```
} ← COEND
```

$\langle \text{name} \rangle * \text{obj} = \text{new} \langle \text{name} \rangle$ ^{SMART}
 $\text{delete obj;} \leftarrow \text{WAIT}$

Explicit concurrency!

④ Actor



#include <uActor.h>

```
struct StrMsg : public uActor::Message {
```

```
    ...
```

```
StrMsg( . . . ) : Message (uActor::Delete / Nodette) { . . . }
```

```
}
```

delete after message consumed
No delete after use

```
- Actor <Actor Name> {
    Allocation receive (Message & msg) {
        Case (StrMsg, msg) {
            msg->msg-d
            if (msg->msg-d == StrMsg)
                Can (BoolMsg, msg) {
                    msg->msg-d
                    return Delete()
                }
            else
                return Nodette()
        }
    }
}
```

Figure out type of msg from denied msg struct.

```
int main ( - - ) {
    uActor::start ();
    * new <Actor Name> | *new StrMsg ( - - ) | . . . | - - | uActor::stopMsg
    :
    will wait until uActor::stop() is called
    uActor::stop ();
}
```

default msg

Implicit concurrency

Actions of _Actor finish:

① Run destructor

② Delete memory

③ Update storage info and remove from actor sys.

Allocation types:

Nodelete \Rightarrow do no ①, ② or ③

Delete \Rightarrow ①, ②, ③ (use for dynamic alloc.)

Destroy \Rightarrow do ①, ③

Finished \Rightarrow do ③. Use for stack allocs

Exceptions

Unhandled exceptions go to last joiner

Same syntax, but faulting exec. not halted as soon as exec. arrives. Checks a

① - Enable begin & end

② suspend / resume

③ yield call

④ -Accept call

Syncing & Comms During Execution

Synchronization: thread waits for another thread to reach some point

\hookrightarrow If sync'd & both threads ready to recv & transmit info \rightarrow comm.

Busy waiting: thread in loop until event happens

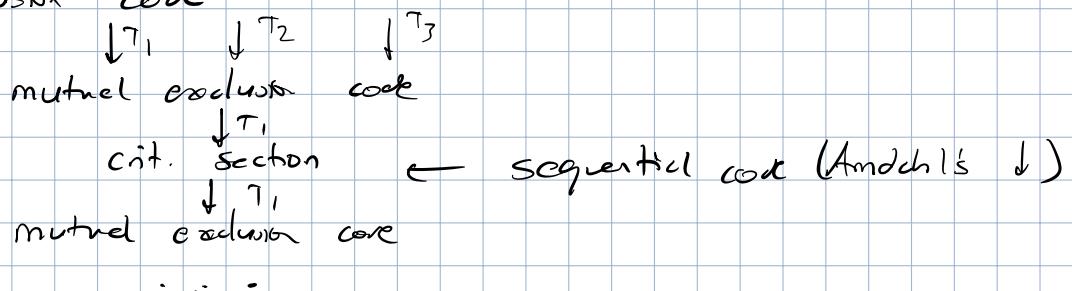
Comm.: changing shared memory (by reference), channels, ...

Critical section

Obj.: want to access an object one thread at a time

L. Note: thread can be interrupted on obj. access, but no other thread can change @ same time

Ans: mutual exclusion code



Bonus: some objects can have multiple threads reading an object \Rightarrow keep read concurrent

Note: static variables are shared in all object instances \rightarrow not safe

Principles of mutual exclusion

Q: Can we write code that forces serial execution?

A: Yes!

Rules that we try to follow:

1. Safety: 1 thread access C.S. @ a time
2. Eventual progress: each thread will eventually execute & can run in any order
3. No selfishness: thread not in entry/exit cannot prevent other code from accessing C.S.
4. No livelock: cannot postpone selection of thread for enter C.S. } Related by
5. No starvation: thread must eventually enter C.S. } busy waiting

Mutual exclusion $\not\Rightarrow$ FIFO service (this is unfairness)

\hookrightarrow Banking: waiting thread overtaken by arriving thread

Software Solutions

① Dekker's Algorithm

```
for (;;) {
    me = Want In
    if (you == DontWantIn) break
    if (Last == me) {  $\Rightarrow$  Forces alternation if went to go in again
        me = DontWantIn
        while (Last == me && you == WantIn) {}
    }
}
CriticalSection()
if (Last != me) {
    Last = me
}
me = DontWantIn
```

Issue: not R/W safe w/out highlighted code (Heselink)

\hookrightarrow Not all R & W are atomic. R/W safe \rightarrow works for non-atomic R/W

Failure Case:

T_0

$Locat = me$
 $me = DontWantIn$
 $(flicker DontWantIn)$

flicker WantIn

flicker DontWantIn
terminates.

T_1

$you == DontWantIn \rightarrow true$
Critical section
Last = me

$you == DontWantIn \rightarrow false$
Wait

↓

Unbounded overtaking: thread can overtake multiple (unbounded #) threads to execute

↳ Thread can come back & re-enter C.S. This is because Last retracts int

② Peterson's Algorithm

```

me = WantIn
Last = me
while (you == WantIn && Last == me) {}
critical section()
me = DontWantIn
    
```

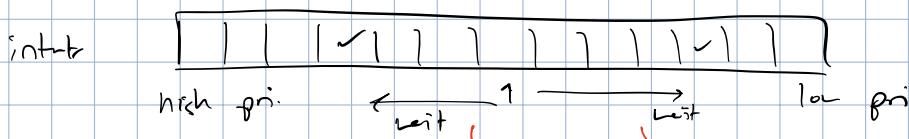


R/W unsafe b/c it needs atomic read & write

Bounded overtaking: loss of Last=me still retains in (doesn't retract int)

↳ Thread cannot come back & re-enter C.S.

③ N-thread prioritized entry



Do {

```

intnts [priority] = WantIn
for (lower indices j) {
    if (intnts [j] == WantIn)
        intnts [priority] = DontWantIn
        while (intnts [j] == WantIn) {}?
    break;
}
    
```

} while (intnts [priority] == DontWantIn)

```

for (higher indices j)
    if (intnts [j] == WantIn) {}?
}
    
```

Check if hi' pri.
Thread has placed
intnt → rechart
& wait

Check lo pri. If wait in → wait

Critical section ()

intents [priority] = Don't wait In

Can have starvation, only need N bits ($N = \# \text{ of threads}$)

↳ Can do 3-bit, but R/W safe 4-bit soln. \Rightarrow R/W safe.

④ N-thread Bakery

Ticket values $\in (0, \infty)$. Lower tickets \Rightarrow higher priority

↳ 0: thread is selected. ∞ : don't wait in

Algo:

```

ticket [priority] = 0
int max = 0
get_max_ticket in ticket arr
ticket [priority] = max + 1
for (int j : 0 → n) {
    while ticket [j] < max || (ticket [j] == max && j < priority)
}
critical section ()
ticket [priority] = INT_MAX
  
```

↑ prevent other threads from entering

} Ticket selection

} Wait for turn.

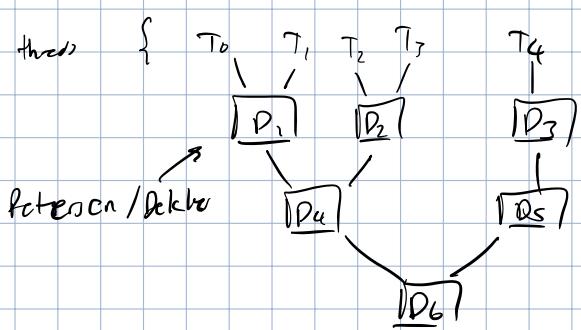
Ticket border

Spec: NM bits
of threads. ticket size

No starvation

R/W unsafe

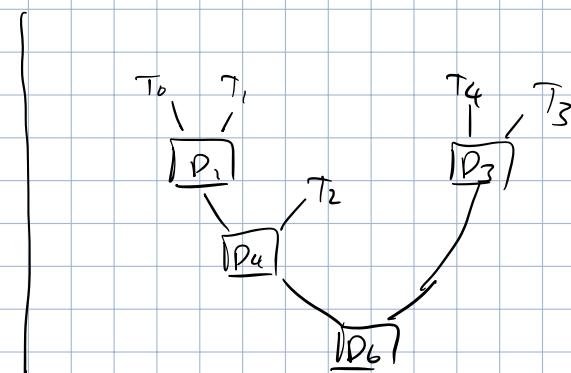
⑤ Tournament



maximci

↓

ensure each thread has same
of D runs



minimci

↓

more unfair → faster

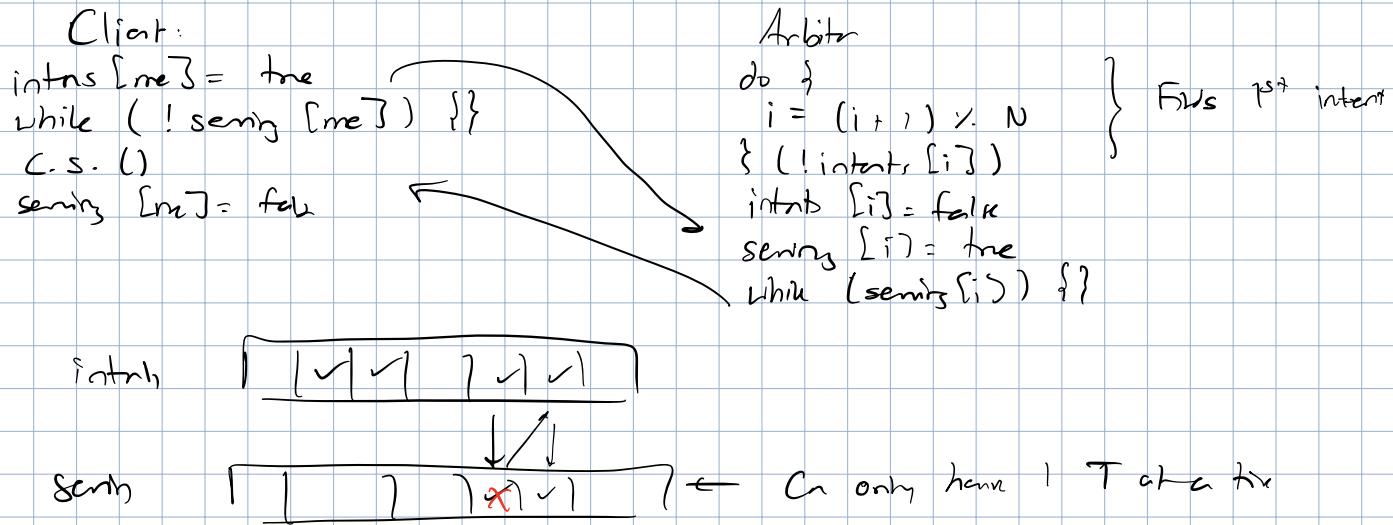
Intents for threads must reverse bottom up

Properties of $D_i \rightarrow$ entire tree has property if all nodes implement it

All have unbounded overtaking

$$\text{Size (minimal)} = \underbrace{(N-1)}_{\# \text{ of tree nodes}} M \begin{matrix} \text{bit} \\ \downarrow \\ \text{node size} \end{matrix}$$

⑥ Arbitrator



Mutual exclusion \rightarrow synchronization

No starvation

R/W unsafe

Con: arbiter thres overhead.

Hardware solutions

SW solutions are complex \rightarrow use hardware.

HW provides atomic read & write ops. \rightarrow cheats!

Atomic instr. :

① Test/Set instr

```
int TestSet (int & b) {
    int temp = b
    b = CLOSED
    return temp
}
```

Coe:

```
while (TestSet (Lock) == Closed) {}
```

C.S.

Lock = Open

Breaker rule S (starvation possible)

② Swap

```
void Swap (int & a, int & b) {
    int temp;
    temp = a
    a = b
    b = temp
}
```



Code:

```
int dummy = Close, lock = Open
do {
    Swap (Lock, dummy)
} while (dummy == Close)
C.S.
lock = open;
```

③ Fetch & increment

```
int FetchIn (int & val) {
    temp = val
    val += 1
    return temp;
}
```

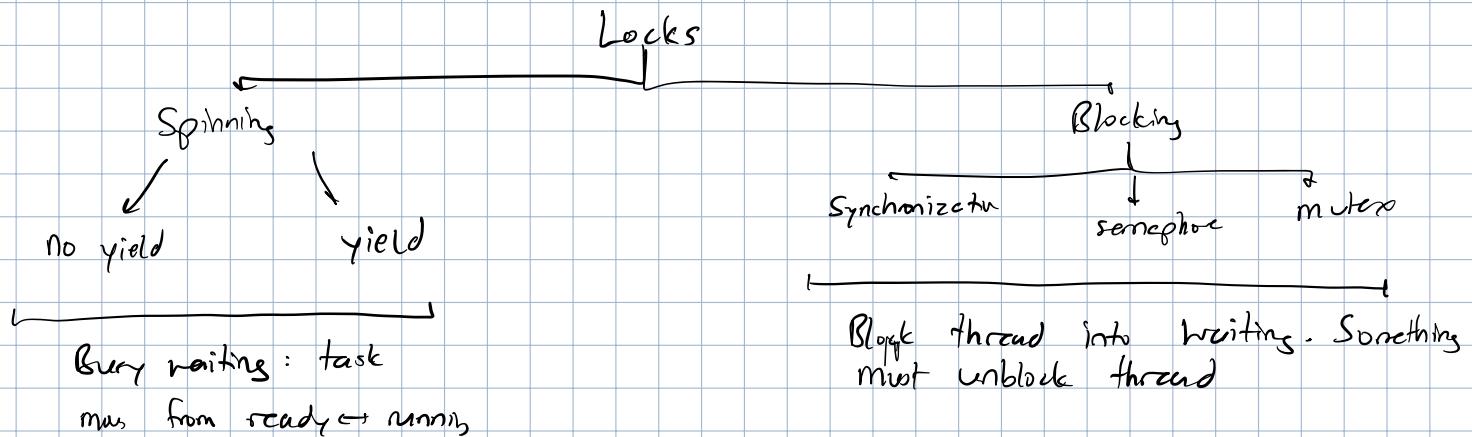


Code:

```
while (FetchIn (Lock) != 0) {}
C.S.
Lock = 0
```

Can overflow, but easy to resolve.

Locks



Spin Locks

Via HW ops: `While (TestSet (Lock) == Close) {}`

Issue: wasteful of CPU cycles

↳ Optimize: yield during spin or move back thread to ready state

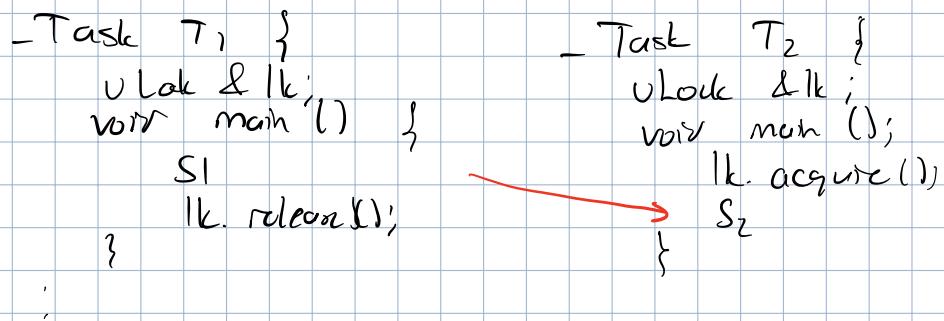
Breaks rule S → starvation possible

Necessary if computer has nothing else to do

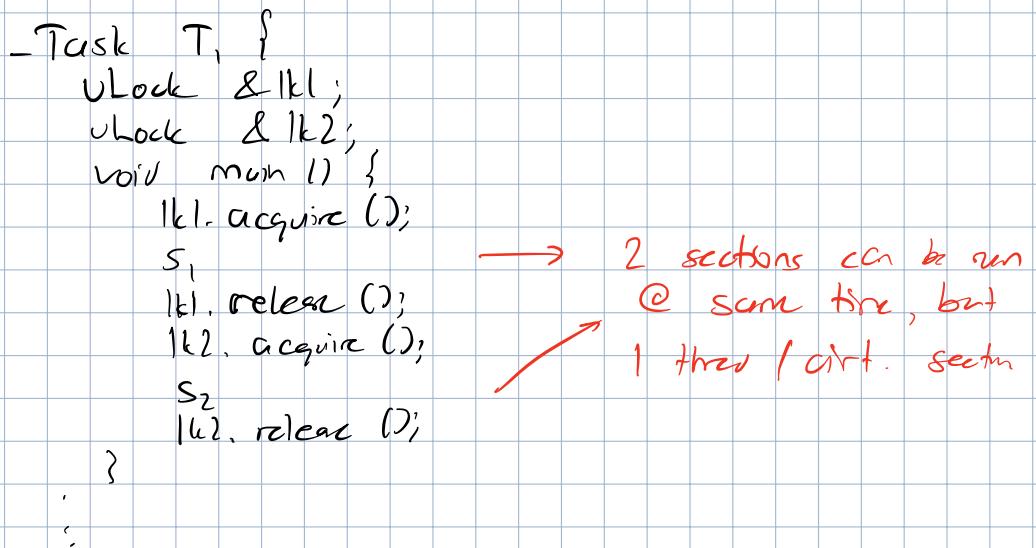
Implementation in C++:

- | | | | |
|---|---------------------------|---|---------------------|
| ① | v SpinLock : non-yielding | } | lk. acquire();
; |
| ② | v Lock : yields | | |

Ex:// Synchronization



Ex:// Multiple locks:



of locks \leq # of critical sections

Blocking Locks

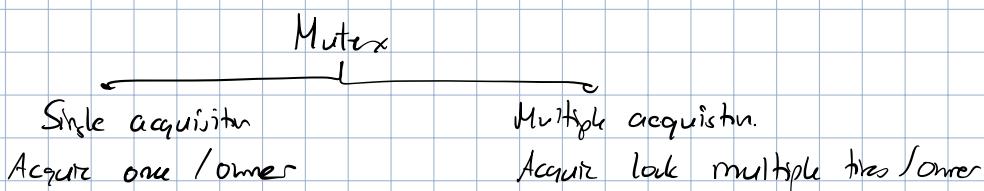
Blocking lock needs someone to release blocked threads → less work, more coordination

All blocking locks have:

- ① State to facilitate lock semantics
 - ② List of blocked acquirers } → any one can be chosen to proceed

A: Mutex lock

Obj: mutual exclusion



Implementation:

```

class MutexClass {
    ULock spinlock;
    Task* owner;
    bool avail;

public:
    void acquire() {
        lock.acquire();
        while (!avail && owner != currThred) {
            yield(&lock);
            lock.acquire();
        }
        avail = False;
        C.S.
        lock.release();
    }
}

```

Basis? thread can come in b/f lock acquire }

2 methods of dealing w/ bargains:

① Avoidance

```
void acquire () {  
    lock.acquire();  
    if (!avail && owner != currThread) {  
        yield (&lock);  
    } else {  
        avail = False;  
        lock.release();  
    }  
    owner = currThread;  
    lock.release();  
}
```

② Prevention

```
void acquire () {  
    lock.acquire ();  
    if (!avail && owner  
        yield (lock);  
    }  
    avail = true;  
    owner = curThread ();  
    lock.release ();  
}
```

```

void release() {
    lock.acquire();
    owner = nullptr;
    if (!blocked.empty()) {
        remove task & make current
    }
    avail = true;
    lock.release();
}

void release() {
    lock.acquire();
    owner = nullptr;
    if (!blocked.empty()) {
        make task ready
    } else {
        avail = true
    }
}
lock.release()

```

↑ pushes control to Blocked thread

μ ++ implementation (multi-acquire): `uOwnLock();`

↳ Must call `release()` == # of times `acquire()` called.

Common pattern:

① - Finally:

```
lock.acquire();
try {
    ...
} - Finally {
    lock.release();
}
```

② RAII

On ctor: acquire lock
on dtor: release lock

Stream locks: osacquire (stream) << "..." << endl; ^ release lock on stream

B. Synchronization Locks

Also known as cond. lock.

Special note: signals to release can be lost if no task present & always block on acquire

Suffers from some busy problem

① Avoidance:

```
milk.acquire();
if (!slk.empty()) s.release();
else { occupied = false; }
milk.release();
```

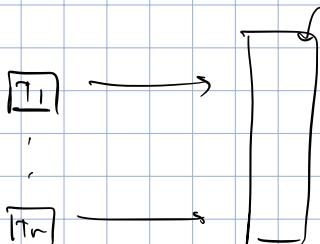
② Prevention.

```
milk.acquire();
if (!slk.empty()) s.release();
else { occupied = false; milk.release(); }
```

μ++ class: CondLock: wait, signal, broadcast
↓ ↓ ↓
block wake up wake up
task all tasks

Need mutual exclusion lock to prevent signal from being lost

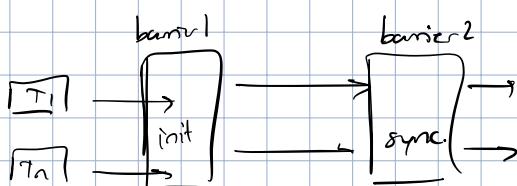
C: Barrier



Barrier forces
threads to wait until all threads good
↓
Sync.

Last task unblocks all tasks on barrier

Another common use:



μ ++ impl. : - Cmonitor

- Cmonitor <Name> : public vBarier {

implicitly
mutually
exclusive

protected:

void last() ← func. executes when last thread enters barrier.

public

void block() ← func. executes if thread calls block()

Cannot use this in COFOR: implicit concurrency → don't know how many threads will be created & need to be in barrier.

D: Semaphore

Equivalent to yielding spin-lock, but blocks instead

Semaphore

Binary

- 0 & 1

Counting

- Multiple states

Common:

- . P(): decrements semaphore value & waits if value is 0
- . V(): increments semaphore value & unblocks waiting task. (usually longest waiters)

① Binary:

Mutual exclusion:

vSemaphore milk(1); initial state

:
milk. P();

C.S.

milk. V();

:

Synchronization:

vSemaphore lk(0);

void f() {

S1

lk. V();

}

void s() {

lk. P();

S2

}

② Counting Semaphores

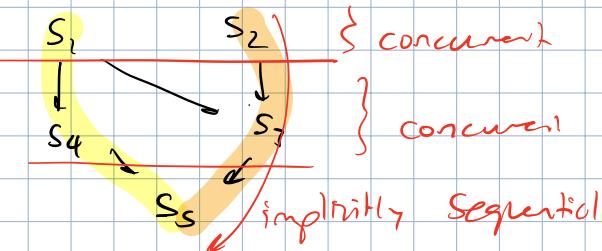
Multiple threads enter critical section → places limit

Lock Programming

Precedence Graph

COBEGIN + Semaphore is as powerful as START/WAIT

Ex: // Code up precedence graph (dependency graph)



Sdn:

v Semaphore L1(0), L2(0);

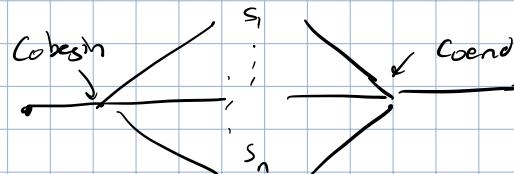
COBEGIN

BEGIN S1; V(L1); S4; V(L2); END

BEGIN S2; P(L1); S3; P(L2); S5; END

COEND

* Note: precedence graph ≠ process graph! Process graph will always look like latter.



Tip: location of P() & V() matters! Only say V() after all state is changed

Split binary semaphores

Defn: collection of binary Semaphores whose sum ≤ 1

Use: different types of tasks have to block separately & need better - partition

↳ Better partition: pass off mutual exclusion to one type of task.
(basing partition)

Also avoid temporal busing (tasks acquiring resource b4 they should).

↳ Sdn: use a queue or private semaphores (each task has own semaphore)

INDIRECT COMMUNICATION

Problem: semaphores require complex prog. We want higher-level concurrency constructs

Critical Regions

Idea: declare variables as shared & guarantee M-E. when changing variables.

```

VAR <variables> : SHARED <types> ← Shared variables that compiler
; ; will guarantee M.E.
; ;

REGION <variable> DO } Provides M.E. for variable.
; ;

END REGION

```

Problems:

- ① You can still read while writing → problems.
 - ② Nested critical regions → deadlock

Conditional region:

REGION <variable> DO
 AWAIT conditional - expr.
;
END REGION .

Monitor

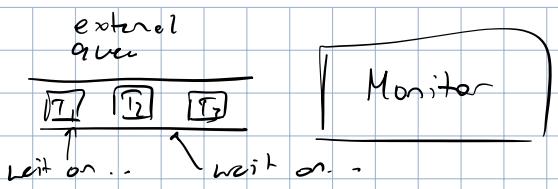
- Monitor <class-name> {
 shared variables => Not guaranteed to have M.E.
 public:
 member1();
 ;
 }
 Public methods guarantee M.E. unless specifl.
- Nonvirtual member k(); => Not M.E. method

Recursive entries allow (multiple acq. bkt)

Dtor releases implicit mutex, so thread will block if in Monitor

Q: How to schedule tasks & provide synchronization?

① External scheduling.



① Internal scheduling

Use condition variables: `vCondition variables; <variables> [..];`

↳ Queue of waiting tasks! Like cond lock

Usage:

```

- Monitor Example {
    vCondition a;
    public:
        void method1 () {
            a. wait();      → Atomically places thread at back of a's
            ;               queue.
        }
        a. signal();   → Remove front of queue & set it ready
    }               ↳ Not allowing signaller to run until
                      signaller finishes method
  
```

`Signal Block ()`: signals thread & blocks signaller

*Important: guarantees no barging! *

When to use external vs. internal?

External should be used if:

- A: scheduler depends on param.
- B: need to block but cannot guarantee signalling

Monitor exceptions

```

void mem1 () {
    ...
    - Throw E();
}
  
```

```

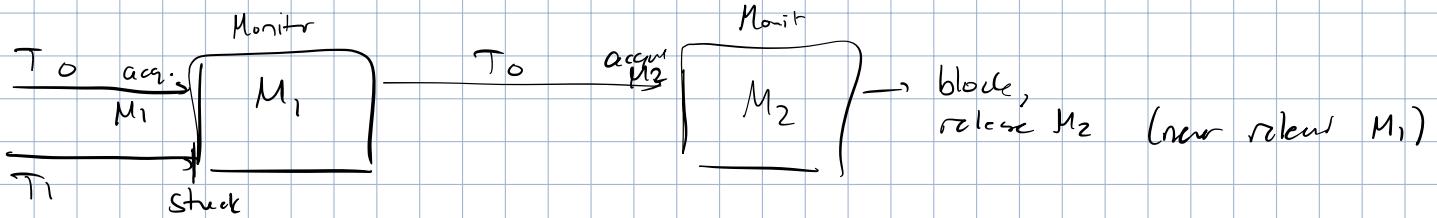
void mem2 () {
    try {
        ...
  
```

```

        - Accept (mel);
    } catch (UMutexFailure::RendezvousFailure &) { . . . }
  
```

Exception type of -Accept methods

Nested monitor



Solution: nest M_2 defn. in monitor M_1 ,

```

-Monitor M1 {
    -Monitor M2 {
        ...
    }
    public:
    ...
}

```

Intrinsic lists

Allocating on heap is a huge bottleneck

↳ Soln: allocate on global stack/heap w/out copying overhead

μ C++ intrinsic lists etc:

- ① `vectorable`: 1-way linked list
- ② `vectorable`: 2-way linked list

{ have their own iterators.

These intrinsic D.S. auto deallocate @ end of scope \rightarrow no dynamic alloc necessary.

Counting semaphore vs. Monitor

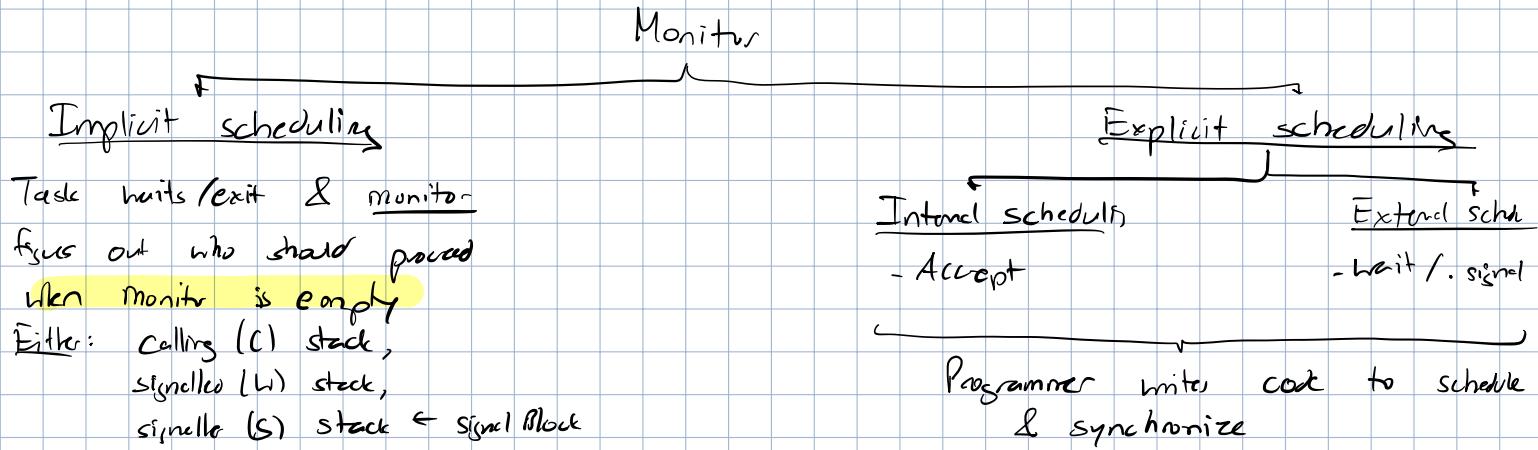
Monitors can simulate semaphores

```

-Monitor Semaphore {
    int sem;
    Condition semCond;
public:
    Semaphore (int cnt=1) : sem(cnt) {}
    void P() {
        if (sem == 0) semCond.wait();
        sem--;
    }
    void V() {
        sem++; semCond.signal();
    }
}

```

Monitor types



We classify monitors based on implicit scheduling priorities

Implicit signalling:

Monitor that waits on predicate & wakes up all waiting threads on mutex member exit to check if pred is $T \rightarrow$ proceed.

```
- Monitor A {  
    - -  
    public:  
        void mem() {  
            waitUntil <pred>;  
            :  
        }  
        void foo() {  
            - -  
        }  
    }  
}
```

- If thread completes:
- ① Wake up all waiting threads
 - ② Evaluate <pred>
 - ③ Continue thread if <pred> = T

Immediate-return signal

```
- Monitor A {  
    - -  
    public:  
        void foo() {  
            - -  
            cond. signal();  
        }  
    }  
}
```

Powerful but not powerful enough for all cases! May need to do more work after signal()

Monitor type chart

signal type	priority	no priority
blocking	$C < S < W$ (signal Block)	$C = S < W \Rightarrow$ while round signal f- benging check
non blocking	$C < W < S$ (signal)	$C = W < S$ (Java/C#)
implicit signal	$C < W$	$C = W$ ↑ while round signalled wait() for benging check

q Good for prototypes, but bad perf.

μ C++ has -C or monitor: coroutine \leftrightarrow monitor

Java Monitors

Use 'synchronized' word to make class a monitor

```
class Foo {
    .
    .
    public synchronized <types> (name) {} . . }
```

Limitations:

- ① 2 factors: `wait()`, `notifyAll()` \rightarrow could lead to benging
- ② 1 condition queue \rightarrow makes certain problems impossible
- ③ Spurious wakeups: waiting threads can randomly wake up & so check if allowed \rightarrow never to have while loop around `wait()`

Sdn: ticketing for benging control & while loops

Also suffers from nested monitor problem

Concurrent Errors

Race condition

\rightarrow poor M.E.

\rightarrow poor synchronization

} threads run assuming in correct state.

No progress

① Livelock

Indefinite postponement on simultaneous arrival

Can always break livelock w/ some tie breaker

② Starvation

A task never executes b/c another stream of tasks always beats it

Long-term starvation is rare

③ Deadlock:

Processes waiting for event that will never occur. Synchronization / M.E.

5 conditions:

A : concrete resource requires M.E.

B : process holds resource & waits for access held by someone else

C : No pre-emption of resource (cannot take away resource)

D : Circular wait

E : simultaneous

Deadlock Prevention

Synchronization deadlock → impossible to statically prevent

Mutual exclusion deadlock:

1. No M.E. → impossible

2. No hold & wait: only give all resources → poor resource utilization

3. Allow pre-emption of resources → cannot do this statically

4. No circular wait: ordered resource policy

① Divide resources into classes

② Only request resource from class R_i ; if not holding resource for any class R_j ($j \geq i$)



5. Prevent simultaneous occurrence

Deadlock Avoidance

Obj: dynamically detect if deadlock occurs \rightarrow prevent situation

Banks' Algo

Check's if safe sequence of allocs \rightarrow no deadlock. Run algo on each alloc.

① Define total amount of resources in sys.

R_1	R_2	R_3	R_4	}	TC
6	12	4	2		

② Define max. # of resources needed / thraw.

	R_1	R_2	R_3	R_4	}	M
T_1	4	10	1	1		
T_2	2	4	1	1		
T_3	5	9	0	7		

③ Define current allocated resources.

	R_1	R_2	R_3	R_4	}	C
T_1	2	5	1	0		
T_2	1	2	1	0		
T_3	1	2	0	0		

④ On resource request, update C w/ request. Calculate available resources ($TC - C$)

$$C_{1,1} : 2 \rightarrow 3$$

$C_B - C =$	R_1	R_2	R_3	R_4	}	CR
	1	3	2	2		

⑤ For each thread \rightarrow figure out how many more resources needed to max ($N = M - C$)

	R_1	R_2	R_3	R_4	}	N
T_1	1	5	0	1		
T_2	1	2	0	2		
T_3	4	7	0	1		

⑥ Choose thread order s.t. all CR resources can be allocated & fulfill some thread's max value (i.e. $CR - N_i = 0$)

Choose $T_2 \Rightarrow$ alloc all resources would make $N_2 = 0$ ($CR - N_2 \geq 0$)

⑦ Give back that thread's resource to CR

	R_1	R_2	R_3	R_4
CR:	2	5	3	2

Repeat until $CR = TC$ (all threads could be run to max & give back)

Allocation graphs

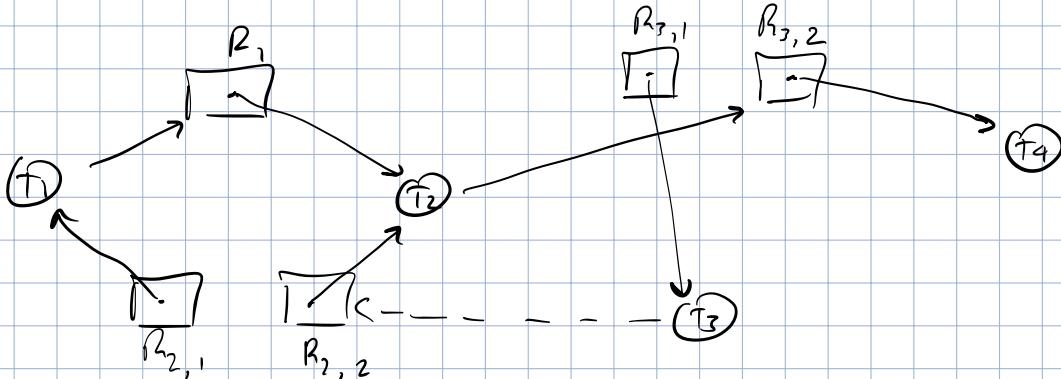
Legend:



Create graph of allocation. No cycle \Rightarrow no deadlock. If multiple instances of resource, then cycle \Rightarrow deadlock

↳ Create isomorphic graphs where 1 node = 1 thread / 1 resource inst.
then cycle \Rightarrow deadlock

Ex: //



Create resource allocation graph on alloc attempts & simulate a path to no deadlock:

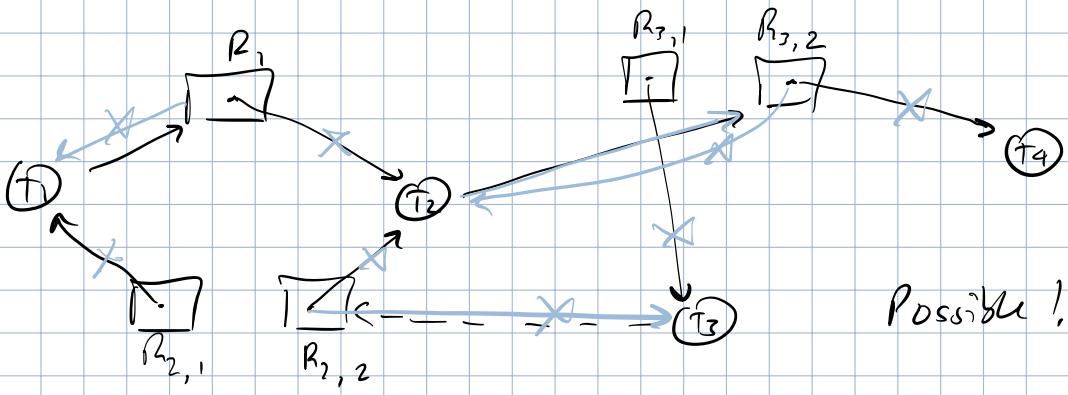
① Create isomorphic

② Thread that is not in hold & wait (e.g. T_4). Release resource.

③ Give resource to T_4

④ Repeat \rightarrow remove all arrows. If possible \rightarrow no deadlock

Ex. 11

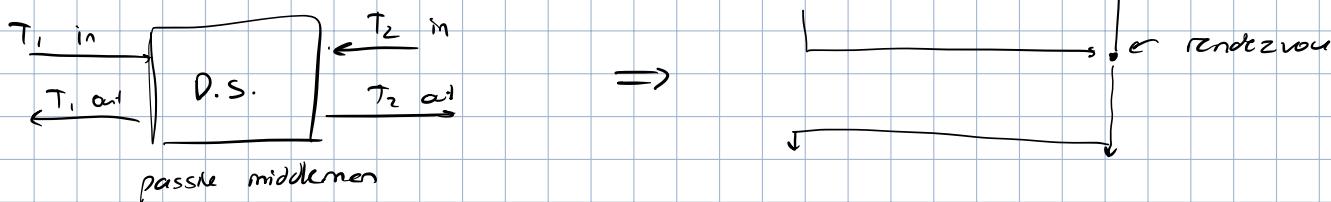


Deadlock Detection & Recovery

Let deadlock happen \rightarrow recover. Need to build allocation graphs & check cycles every T sec. Need preemption

DIRECT COMMUNICATION

Objective:



* Key: a -Task is like a -Monitor \rightarrow only 1 thread can be active in -Task @ a time

↳ Public members already have -Mutex

↳ If T1 calls -Task public member \rightarrow -Task can block & allow caller to execute via external sched

How to make a Task into monitor:

- Task A {

private:

public:

```
void foo() { .. }  
void bar() { .. }
```

private:

```
void main() {  
    for(;;) {
```

Should not need to -Accept. Make this as small as possible \rightarrow -Task that does work

gets blocked \rightarrow other threads can call public method

When (<cond>) -Accept (foo) { .. }
or When (<cond>) -Accept (bar) { .. }

ran by -Task after call done

- Task Thread

Conditional -Accept

Notes:

- ① If multiple calls accept & outstanding call to call (e.g. -Accept (foo & bar))
 - ↳ Execute {..} of earliest -Accept → order matters
- ② {..} executes in 2 situations:
 - A: function accepting finishes
 - B: earlier calls wait() in function that re-Accept
- ③ If using interval scheduling, -Task thread must call signal Block()
- ④ To stop monitor, do: -Accept (~className) } return; }
 - ↳ Caller gets blocked while {..} executes → ~className() porters

Increasing Concurrency

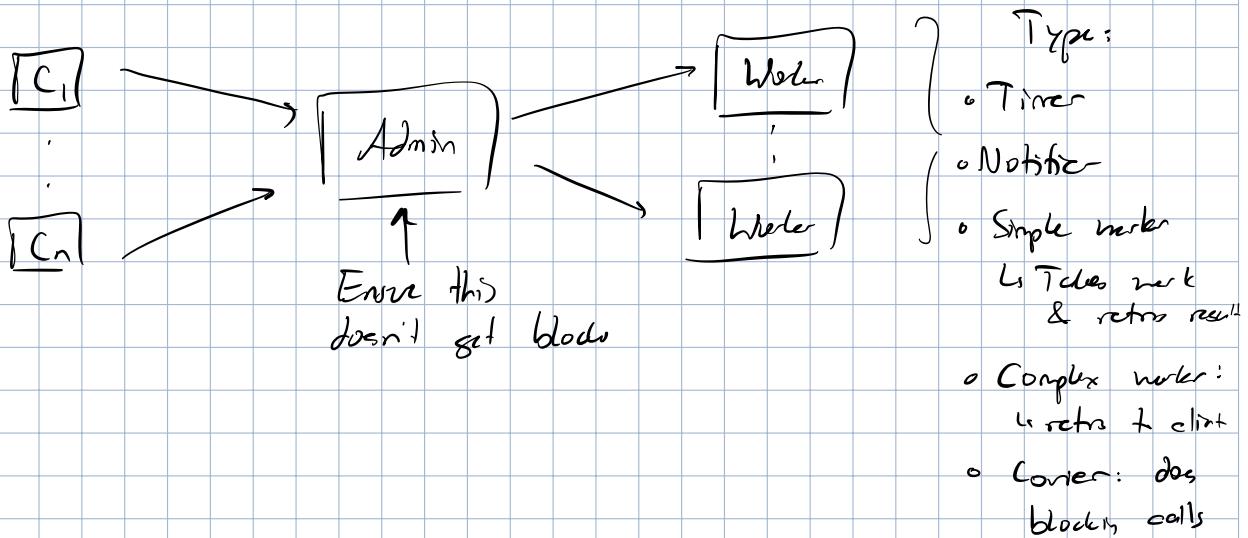


Q: How to increase concurrency on server-side?

Default -Task method has a problem:

- ① M.E. of function ⇒ client will get blocked waiting for server to do work

Idea: separate out admin from work



Notes:

- ① Admin can have branch for workers to wait while work is getting ready
- ② Admin will never have deadlock
- ③ Client will drop off work to admin in internal buffer

Q: How to increase concurrency on client-side?

Obj: asynchronous calls + server → client shouldn't wait.

↳ In µ++: no async calls → build async. out of sync. calls

Methods:

① Value returns

Separate out call & return

```
server.call(); } can be done asyn. to  
work { serv  
serv.wait(); ← could block
```

Issue: server needs to know which client is waiting & give proper result

② Tickets

Do same thing but server gives client ticket to identify client

Issues: assuming no malice from client (e.g. ticket is false final)

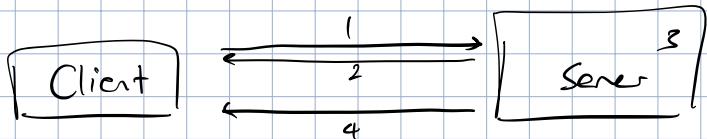
③ Call-back routines

Transmit routine on call → serv will call routine on work completion

Assumes call-back is not blocking

Client must wait/poll until call-back called, "push" result from serv

④ Futures



- 1 : client asks for result
- 2 : server immediately gives future
- 3 : server computes result
- 4 : server delivers result to future in client

If client wants to use future, must wait for step 4.

General design:

Future maps result

```

class Future<public ResultType> {
    friend Task serve; → server can check future (step 4)
    ResultType result;
    vSemaphore avail;
    Future* link; → linked list of futures for server to fill in
public:
    ctr. ~
    ResultType get() {
        avail.P(); → Caller blocks here. Server must signal that
        return result; result is ready
    }
}

```

M++ futures:

Client:

```

#include <UFuture.h>
.

Future<ISM<ResultType>> foo; } Declaration & get
foo = serv.perform(...); } future
.

ResultType result = foo(); → Wait on future. This must
; be done!
foo.reset(); → make future empty & restart
;
foo.cancel(); → Stops server & client from using it
;
foo.available(); → T if async call complete, F o.w.

```

Pointers of futures cannot change

Server

① Keep reference of created futures

② Deliver result to future

int work = ...
future.deliver(work); Assuming work matches ResultType of future.

Can also deliver exception

future.deliver(new Exception());

Future will handle dynamic errors for you

Complex Access

- Select(f1) { ... } or - Select(f2) { ... } and - Select(f3) { ... }
↑
executes on f1 delivery ↑
Both f2 & f3 must
be delivered to make forced
BUT { ... } is executed on singular
delivery

If select sequence not satisfied → block

Conditional select :

- When (...) - Select (...) }
or - When (...) - Select (...) } toggle - Select on & off
and - When (...) - Select (...) }

- Else also available → runs if - Select statement block.

OPTIMIZATIONS

Type of optimizations

Reordering

Data & code reordered
for perf.

Eliding

Removing unnecessary
code & data

Replication

Duplicate across
resources (e.g. concurrent)

Goal: faster progr. but should be isomorphic (input → output is same)

Note: more optimizations → compile time ↑

Sequential Optimizations

① Reorder disjoint operations

{ t = x;
s = y; , , z = 3; }

② Elide output.

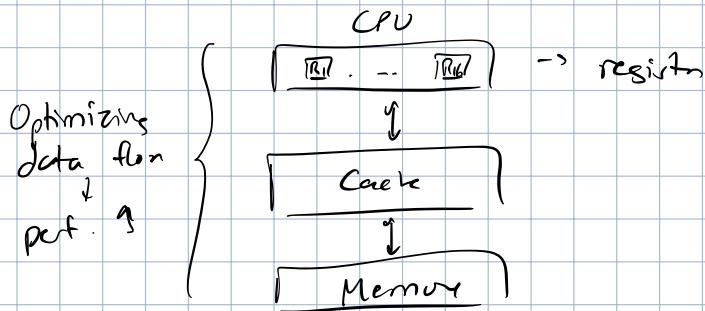
$$\begin{array}{l} x = 0; \\ \cancel{x = 3}, \end{array}$$

~~for (int i=0; i < n; i++)~~;

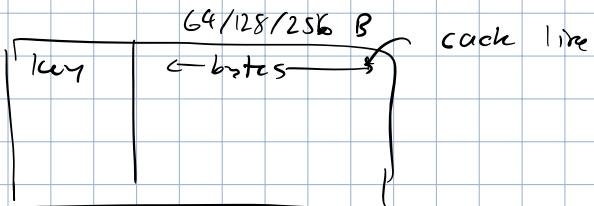
③ Execute code in parallel if multiple units (e.g. put int add & float add in diff adders)

↳ Micro-parallelism in L1.L2., limited in sequential

Memory Hierarchy

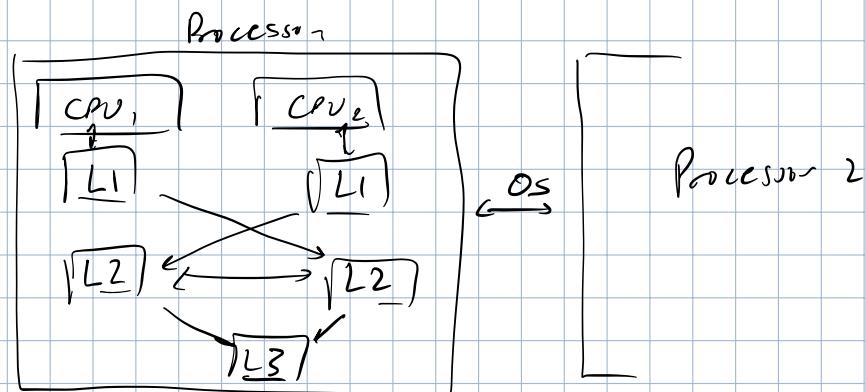


Cache: fast access to mem. that cannot be stored in registers for time



- Preordering: enables faster access to data binned off cache

Cache coherence:



Q: How to reconcile cache diff. across cores / processors?

A: Cache coherence protocol from H.W.

Cache consistency: CPUs see data in an order that makes sense

Consistency + coherence \Rightarrow atomic R/W

Total time \downarrow \Rightarrow threads will read stale values. Block until values updated

Multithreaded problems w/ cache:

- ① Cache thrashing: multiple threads update same cache line → cache update ↑
- ② False sharing: multiple variables, thrash w/ entire cache line → thrashes

Concurrency Optimizations

Sequential prog. have strong memory ordering: value reading is always last value written

Concurrent || weak || : not guaranteed to be reading right value

Concurrent prog.:

concurrent code

Seq. code } optimize, but can still cause errors
conc. code

Examples of seq. optimizations → concurrent errors:

① Disjoint ordering

A: Reordering reads → not a problem

B: Reordering reads b/f writes → problem

C: Reorder writes b/f read → problem

D: Reordering writes → problems

E: Reorder locks acquisition & release → problems

↳ Double-check locking: only allow 1 thread to set ip

int* ip = nullptr;

...

if (ip == nullptr) { ↗
 lock.acquire();

 if (ip == nullptr) { ↗

 ip = new int(...);

 }

 lock.release();

}

Check twice?

2 threads come in, both
see ip == null, both
will try to change
it 2nd not there

This doesn't work! Compiler reorders malloc & setting ip to null.
so multiple threads can still set ip!

① Eliding

Remove code used in other threads → problems.

Memory Model

Manufacturers set model → which optimizations H.W. will perform

↳ Compiler can use this to prevent optimizations from messing concurrent prgs.

Minimum model for concurrency: sequential consistency

Prevent Optimization Problems

Goal: race-free prgs. (produce races if M.E & sync provided)

↳ Doesn't mean no races. Races intend to crit. section (up to prg.)

Approaches:

① Augment code w/ pragmas → easier

② Formally define races & prevent

Tricks:

① volatile: forces load & store of variable

↳ In Java/C++: atomic

② Set fences: bounds compiler reordering

Hard!

OTHER APPROACHES

Atomic / Lock-Free Data Structures

Obj: no ownership over D.S. but still only allow atomic changes

Misnomer ⇒ still spinning. If eventual progress → wait-free D.S.

Based on atomic CAS (compare & set):

```
bool CAS(int& val, int comp, int nval) {
    if (val == comp) {
        val = nval;
        return true;
    }
    return false;
```

} If val == comp, set.
Return if val == comp

{

Ex:// Stack:

```
void Stack::push (Node* n) {
    for (;;) {
        n.next = top;
        if (!CAS (top, n.next, &n)) break;
    }
}
```

Ensures $\text{top} = n$
atomically

```
Node* Stack::pop () {
    Node* t;
    for (;;) {
        t = top;
        if (t == nullptr) return t;
        if (!CAS (top, t, t->next)) break;
    }
}
```

Problem: ABA

$\xrightarrow{\text{top}}$

Stack: A, B, C]

- ① T₁ calls pop, $t = A$. B/T CAS, time sliced
- ② T₂ pops A, B, pushes A
- ③ T₁ returns. Check if $\text{top} == A \rightarrow \text{true!}$ Set top to B (but B already popped!)

Fix:

- ① New atomic inst.

double-wide CAND atomic instr.

```
bool CAND (int& val, int& comp, int nval) {
    if (val == comp) {
        val = nval;
        return true;
    }
    comp = val;
    return false;
}
```

- ② Have tickets for each node on stack

Every push has increased count on data struct. Only pop if right counter.

Only probabilistic situation

Exotic atomic instructions

H.W. provides strong atomicity \rightarrow create more complex lock-free D.S.

Ex:// MIPS has LL to put value into reg., SC to write back
set a reservation \rightarrow tells if reg. overwritten!

On SC \rightarrow error if oracle says reg. is overwritten.

GPU

GPU uses SIMD arch \Rightarrow all instr. on CPUs must occur at some time

Makes prog. difficult, esp. branching \Rightarrow must execute all instr. even if branch = false!

\hookrightarrow Requires complex scheduler to prevent unless ops + speedups

No cache, only registers.

Great for vectorized ops.

Other Languages

Ada

- ① When needs to be at top & can only use class variables
- ② No internal scheduling (all when, implicit signalling)
- ③ Problem w/ blocking (reqd) after computation \rightarrow lose everything!

Java

- ① Requires start() & join() to be called

Go

- ① Non-pre-emptive user threads
- ② Channel comm. vs. Object comm. in C++

C++

- ① Call:
`thread objName (routine, params);`
- ② Need to call join()
- ③ Has many concurrency constructs but monitor is missing

java.util.concurrent

Multiple condition locks!

threads